



Biofilm classification analysis pipeline User manual

Version 1.0.0

IMAGINE Project

October 3, 2025

Contents

1	Introduction	2
2	Image requirements	2
3	PC Requirements	5
3.1	Docker installation	5
3.2	Building the Docker image	5
4	Preparation of data for running the classification pipeline	6
5	Generating features using the classification pipeline	7
6	Classification using the pipeline	9
6.1	Explanations of individual files	11
6.1.1	Classification_all and classification_no_counts_features	11
6.1.2	Ablation_ranking	12
6.1.3	Rankings_label and rankings_date	12
7	Visualisations of the comparisons of the performance of classifiers and the ablation ranking	13
8	Examples of applications of the developed pipeline to optimise the experimental setup or to use the pipeline to predict other properties related to biofilm characteristics	13
8.1	How do you use the pipeline if you want to link biofilm features to properties other than the strain name?	13
8.2	Using the pipeline to determine whether the number of layers in a 3D image affects classification performance	14
9	Feature descriptions to help with interpretation	14
9.1	Morphological features	15
9.2	The statistics-based intensity features	16
9.3	Texture features	17
9.4	Large language model (LLM) based features (GPT features) .	18
A	List and descriptions of biofilm features calculated by the pipeline	19

1 Introduction

The biofilm analysis pipeline automatically extracts features from biofilm images, identifies patterns in three-dimensional data and classifies images based on a selected target. Here we present the use when the target is the name of the strain and the classification process is used to determine which bacterial strain each image represents.

2 Image requirements

As input we use 3D images (z-stacks) of biofilms consisting of several image layers. They must be saved in tif format, other formats are currently not compatible with the pipeline.

In case of Zeiss LSM10 microscope high-throughput imaging, the images at different positions of the sample can be acquired in such a way that the images are saved in one file. The images for further analysis must be separated into separated tiffs (one image represents only one position, which includes all images of a 3D stack). This can be done using open source software [ImageJ FIJI](https://imagej.net/software/fiji/)¹ with a provided ImageJ script for splitting and converting images from lsm files (Zeiss microscopes) (Figure 1).

The execution of the pipeline depends heavily on the information contained in the image names, which should follow the example below.

Scheme for image naming:

`<DT>_s_<SP>_st_<ST>_p_<WC>_pos<WP>_tm_<GT>_ch_<CH>_z_<LD>`

- **DT** - date
- **SP** - species
- **ST** - strain
- **WC** - well coordinates
- **WP** - well position
- **GT** - growth time
- **CH** - channel

¹<https://imagej.net/software/fiji/>

- **LD** - layer distance

Example image name:

07112023_s_Lm_st_L1823_p_F02_pos005_tm_24_ch_Syto9_z_21

Explanation of the name:

- 07112023: date in DDMMYYYY format
- s_Lm: s_ indicates the species, Lm means *Listeria monocytogenes*
- st_L1823: st_ denotes the strain, L1823 is the name of the strain
- p_F02: p_ denotes the well coordinates on the 96-well plate, F02 is the actual position on the 96-well plate
- pos005: pos denotes the measured position within the individual well, e.g. position 5 in well F02
- tm_24: tm_ stands for the measurement time, 24 means that the images were taken after 24 hours of biofilm growth
- ch_Syto9: ch_ denotes the channel, Syto9 is the name of the dye and filter used
- z_21: z_ denotes the number of layers in the 3D image, which also reflects the distance between two consecutive layers.

The image names are parsed on the back-end. Thus, following the naming structure is very important. Any deviation from the described pattern will result in processing failure.

The renaming process can be carried out semi-automatically by the [software](#)², or by user-defined scripts. Examples and the user manual are available [here](#)³ together with graphical representations.

²<https://www.advancedrenamer.com/>

³https://www.advancedrenamer.com/user_guide/v3/complete_guide




Figure 1: The script enables the conversion of an lsm file with several 3D stacks taken in different wells into single 3D stack per position (one-position-one-set of images – one 3D stack) in tif format. The script is loaded by dragging it onto the ImageJ; once it opens, the number of images per well and the number of wells in which images are captured must be adjusted. This script also relies on the image naming shown below. If you do not follow these, the script will not work.

3 PC Requirements

The pipeline is designed to run on both Windows and Unix-based systems, or more generally on any system that supports Docker. Hardware requirements include at least two computing cores (CPUs) and 4 gigabytes of ram. The files required to run the pipeline are approximately 500 MB in size, and 1 GB of memory may be required to run the analysis smoothly.

Level required:

- User level of windows environment navigation. Acquaintance with command line in Windows, basics available [here](#)⁴.
- Acquaintance with Docker environment, basics available [here](#)⁵.

3.1 Docker installation

A Docker host daemon is required to build Docker images. The easiest way to meet this requirement is to install the Docker Desktop application.

1. Download and install the Docker Desktop application from the website. This might take a minute, depending on your system and will require you to restart your computer. [Link Windows](#)⁶ | [Link Ubuntu Linux](#)⁷.
2. Start Docker Desktop by double-clicking on the Docker icon in the Start menu. It might run in the background and you might need to bring it to the front from the tray (bottom right corner).

3.2 Building the Docker image

The next step is to build the Docker image.

1. Clone the source code from the GitHub website into the desired location.
2. Start the terminal in the above-selected location.
3. Build the biofilm classification Docker container using the command below.

⁴<https://www.geeksforgeeks.org/most-useful-cmd-commands-in-windows/>

⁵<https://www.docker.com/blog/getting-started-with-docker-desktop/>

⁶<https://docs.docker.com/desktop/install/windows-install/>

⁷<https://docs.docker.com/desktop/setup/install/linux/>

For Windows users:

```
c:\my-download-folder\biofilm-classification> build-and-verify.bat
```

For Linux users:

```
# to make the file executable
```

```
(/my-download-folder/biofilm-classification) $ chmod +x build-and-verify.sh
```

```
(/my-download-folder/biofilm-classification) $ ./build-and-verify.sh
```

Once the Docker image is built and verified (it takes a while), you can run analyses as described below. The Docker image will not need to be rebuilt unless changes are made to the source code.

You can close the Docker desktop application but leave it running in the background (you will see an icon in the taskbar indicating that the Docker daemon is still running and ready to run containers). You can close the terminal window if you do not intend to continue or leave it open to run the classification pipeline.

4 Preparation of data for running the classification pipeline

Before executing the command to calculate features you need to:

1. Prepare a folder with your input images, e.g. create a folder `example_images` on the desktop, and remember the path or address of the folder, e.g. `/c/Users/my_name/Desktop/example_images`).
2. Create a folder in which you want to save the calculation results. For example, create the subfolder `results_example` on the desktop and make a note of the path or address, e.g. `/c/Users/my_name/Desktop/biofilm-classification/results_example`.
3. (optional) Change the voxel dimensions if necessary: To do this, go to `src` folder, find the feature generator script and change the following part as desired:
constants `voxel_size = 0.13 * 0.13 * 0.5` `pixel_area = 0.13 ** 2`. If you make changes to the feature generator code, you must recreate the Docker image as described above.
4. Open the source code folder. Search for the file with the name `".env"`.

You will need to replace the path to the images in the .env file:
"/c/Users/my_name/Desktop/example_images" with the path or address where your images are located. You can check this in the address bar of the window when the folder with the images is open. This address needs to be changed slightly to be compatible with Docker by following the rules below: change " to "/" and remove ":" as in the following example:

c:
my_folder
test_images, you need to write it as /c/my_folder/test_images. Make sure that you do not change the text after the ":".

One must replace "/c/Users/my_name/Desktop/biofilm-classification/results_example" with the corresponding path if you have created the subfolder in a different location than described above. The address must be changed as described above. Again, make sure that you do not change the text after the ":".

Open the file and change the paths in the .env file:

```
# change these paths so that they point to your data and desired results folder
IMAGINE_IMAGES=/c/Users/my_name/Desktop/example_images
IMAGINE_RESULTS=/c/Users/my_name/Desktop/biofilm-classification/results_example
```

5 Generating features using the classification pipeline

To generate features with the pipeline, run the following command in the terminal (the parts in<> must be replaced as described above and below and <> as well (see examples below):

```
docker compose run --rm datafile.tsv generate_features
```

If you are not using an .env file but your own file, change the command accordingly:

```
docker compose
  --env-file myownfilename.txt
  run
  --rm biofilm-classification
  datafile.tsv generate_features
```

For the examples above, the command is:


```
docker compose run
    --rm biofilm-classification 1
    datafile.tsv 10 generate_features
```

And, if you use your own file with paths:

```
docker compose
    --env-file test_paths.txt run
    --rm biofilm-classification 1 datafile.tsv 10 generate_features
```

The parameter “1” determines the parallelization rate – a higher parallelization (should not exceed the number of CPU cores in your computer) works faster, but consumes more resources of the computer, while a lower parallelization (1) works slower, but should also be completed on machines with fewer resources.

The parameter “10” determines how many of the top features are used in the visualizations.

If you want to test the software, you can use test images (which are supplied with the software) by executing the test_images folder as:

```
docker compose run --rm imagine-test-generate-features
docker compose run --rm imagine-test-learning-benchmark
docker compose run --rm imagine-test-data-visualization
docker compose run --rm imagine-test-learning-benchmark-save-models
docker compose run --rm imagine-test-inference
docker compose run --rm imagine-test-inference-env
docker compose run --rm imagine-test-inference-fast
```

The main output of this step is **datafile.tsv**, where all calculated features are available for review. Once the process is complete, check the results folder and in particular the datafile. In the first column of the data file, all analyzed images should be listed by name. If some images are missing, check the format of the missing images.

An example of what kind of results you can expect can be found in the biofilm-classification/examples/results_example folder that belongs to the pipeline.

The file datafile.tsv can be opened with Excel for checking. If the file becomes too large for Excel, selected rows of the table can be exported to csv and opened with Excel to check the calculations more intuitively: Open the command prompt as described above. Navigate to the folder with the results

using the `cd` command as described above. Then use the command:

```
type datafile.tsv |  
Select -First 15 > first_15_rows_of_datafile.csv
```

The features from the data file are now prepared for running the classification model.

How features are calculated is described in the last chapter of the manual and in the attached table. Further details can be found in the code itself. Many examples are also given in the manuscript based on the use case of *Listeria* biofilms.

6 Classification using the pipeline

The command for comparison of the classifiers (benchmark) is as follows (the `<>` sections must be replaced as above):

```
docker compose run  
    --rm biofilm-classification  
    datafile.tsv learning_benchmark
```

For example:

```
docker compose run  
    --rm biofilm-classification 1  
    datafile.tsv 10 learning_benchmark
```

The parameter “1” determines the parallelization rate — a higher parallelization (should not exceed the number of CPU cores in your computer) works faster, but consumes more resources of the computer, while a lower parallelization (1) works slower, but should also be completed on machines with fewer resources.

The parameter “10” determines how many of the top features are used in the visualizations.

For the classification comparison, the set of given images is randomly divided by the software into a training set and a prediction set. The training set is used to train the model and the prediction set is used to test how accurately the names of the strains can be predicted based on the trained model. These results show us how successfully the strains could be predicted.

The results of the comparison of the classifiers are contained in three different files:

- classification_all
- classification_no_counts_features
- ablation_ranking_all
- rankings_date
- rankings_label

Label here corresponds to the string that follows `_st_` in the name of the image, in this case strain name. The classification target is the property predicted by the classifier, and in this case the target corresponds to the label, i.e. the name of the strain. The date refers to the exact date on which the experiment took place and is linked to the biological replicas. The date is also indicated in the name of the image (see naming convention).

Under some circumstances (most notably low available memory) some of the models (AutoGluon) can unfortunately crash the container, however, the software is designed in such a way to store intermediary results so most of the computation will not be wasted. If the software exits prematurely, you can try to rerun it and it should pick up where it left off. In case it stops again, try closing other applications that might be using up the available memory. In the case where a model will not complete even after several tries you can disable it using the instructions provided below.

Disabling individual models: The file that determines which models will be used in the benchmark is `src/feature_ranking_lite.py`. If you wish to disable some of the models locate the line 268; starting from there the following lines should include a definition of which models are used, e.g., as such:

```
models = {
    'dummy': DummyClassifier(),
    'decisiontree': tree.DecisionTreeClassifier(),
    'logistic': LogisticRegression(),
    'rf': RandomForestClassifier(),
    'xgb': XGBClassifier(
        n_estimators=100,
        max_depth=3,
        learning_rate=1,
```

```

        objective='binary:logistic'),
    'gridsearch': GridSearchCV(
        KNeighborsClassifier(),
        parameters,
        n_jobs=PARALLELISM),
    'autogluon': TabularPredictor(label="label"),
}

```

Any model can be disabled by placing a hash symbol # immediately before its definition. For example, to disable "autogluon":

```
#'autogluon': TabularPredictor(label="label"),
```

After any such modification in the selected models, the image should be rebuilt by using:

```
docker compose build --no-cache
```

6.1 Explanations of individual files

6.1.1 Classification_all and classification_no_counts_features

These two files summarize the results of classification by selected machine learning algorithms. The columns from left to right are:

- tag – has value RESULT and it has been used for parsing, but is currently not active
- upsampling – has value 1
- model - type of classifier, random forest, etc. and dummy classifier as sanity test
- n_components - number of input features

The classifiers can be fed with all features, currently there are 2700, or the number of features can be reduced beforehand by the SVD decomposition model, which selects the features it considers important and outputs the number of features the user selects - here we reduce the number of features by a factor of two; when comparing the accuracy for the reduced data, you should bear in mind that the features selected by SVD are not the same for all dimensions, which means that the model can be fed with different features and therefore produce different results.

- fold – number of iterations (repetitions)
- accuracy - number of correct predictions (key result of the comparison)
- test_set - name of the strains (labels) used for the classification test,
- thr_features – TRUE if features with threshold are included and FALSE if features with threshold are excluded; for each number of n_components results for TRUE and FALSE are given; for n_components 580 we have only FALSE (all features except the one with threshold) and for 2700 only TRUE (because all features include the one with threshold) – so 580 and 2700 should thus be considered and compared together

For the task of classifying biofilm images by strain name, we choose the classifier random forest. The following files contain information derived exclusively from classifications based on random forest.

6.1.2 Ablation_ranking

The report shows how model accuracy changes as features are added incrementally. Determining the highest accuracy for a dataset shows how many features are required to classify the strains. For interpretation, we recommend selecting the number of features that gives the highest accuracy and then using the same number of features with the highest score to explain which ones contribute the most to the classification. For example, in ablation_ranking we find that the number of features with the highest accuracy is 81; then we go to the ranking file and select 81 features with the highest score for further interpretation.

6.1.3 Rankings_label and rankings_date

Label here corresponds to the string that follows _st_ in the name of the image, in this case strain name, and in the classification we call the property that is predicted, the target here is either the label (strain) or the date.

For each feature calculated, the strength of its relationship to the target is indicated, with larger values indicating that the feature was more important in predicting strain and vice versa (rankings_label). We also provide rankings for date prediction (rankings_date) to account for features that may vary as a consequence of repeating the experiment on different days (likely due to the methodology or variability of the biological system under study). This

is crucial when we provide multiple images from different days (biological replicates) for a single strain to the model, as the variation within a single day is often less than between different days. If the model relies on features describing biofilm characteristics that are tied to specific dates rather than strains across different dates, it may lead to bias and poorer results. We suggest that the features included in `rankings_label` and `rankings_date` are compared and that the high-ranking, overlapping features are not considered when explaining the complexity of the system.

7 Visualisations of the comparisons of the performance of classifiers and the ablation ranking

The results of the performance comparison between the selected classifiers (based on calculations from the `classification_all` and `classification_no_counts_features`) and the ablation ranking (`ablation_ranking`) are now visualised to facilitate the interpretation of the results. These are found in the folder “Visualizations”.

8 Examples of applications of the developed pipeline to optimise the experimental setup or to use the pipeline to predict other properties related to biofilm characteristics

8.1 How do you use the pipeline if you want to link biofilm features to properties other than the strain name?

The easiest way is to replace the string after `_st_` in the image name with the desired property. For example, to associate biofilm features with clinical relevance, replace the string after `_st_` (strain name) with `CI` (clinically relevant) or `CNI` (clinically not relevant) based on the epidemiology of the strain. Then run the pipeline as usual.

8.2 Using the pipeline to determine whether the number of layers in a 3D image affects classification performance

The number of layers in the 3D images can be reduced with the command from the pipeline:

```
docker compose run
    --rm biofilm-classification datafile.tsv
    reduce_layers <initial number of layers>
```

To reduce the layers from 21 to 11, you can use the following example:

```
docker compose run
    --rm biofilm-classification 1 datafile.tsv 10 reduce_layers 21
```

This command removes every other layer from the bottom up to create a 3D image with half of the original layers. We can reduce the number of layers down to the bottom layer. Then use images with a reduced number of layers to run the pipeline as described above.

9 Feature descriptions to help with interpretation

A set of quantitative features is extracted from each 3D image (Appendix A) The 3D images here consist of several layers, e.g. we have a 3D image consisting of 21 layers (Figure 2). Understanding that 3D images consist of several sub-images is crucial for understanding how calculations are performed.

Image analysis usually involves some pre-processing steps to distinguish the background from the foreground and to help to recognise objects in the image. Since these procedures require human input, we try to avoid them. In general, the input for feature calculation is a raw image, but for feature counts the pixel intensity values in the raw image are normalised. Normalisation is done by subtracting each pixel intensity value from the average pixel intensity of that layer and then dividing by the difference between maximum and minimum pixel intensity $(x - \text{average}) / (\text{max} - \text{min})$ to adjust the range of pixel values in an image to a standard range, which here is between 0 and 1. If the feature name in the data file contains “Norminten”, the image has been normalised; if the raw image was taken as input, this is not the case.

All features provide a value for each image but are calculated in two different ways: a) the feature is computed immediately for the whole image or b) the




Figure 2: Schematic representation of input images. Green rods on the left image represent bacteria within the biofilm.

layers are first considered individually and then a value per image is obtained by aggregating the layer values by computing the mean (mean), median (med), standard deviation (std), variance (var) and quantiles (q10, q25, q75 and q90). If the feature is aggregated in the data file, this is noted as `_mean`, `_med`, `_min`, `_max`, `_std`, `_var`, `_q10`, `_q25`, `_q75`, `_q90` at the end of the feature's name. In addition, we also introduce normalised values of features that are calculated for each layer as described later.

Below you will find descriptions of the features. Appendix [A](#) lists the abbreviations from the data file, the full names of the features, the processing steps and brief explanations.

9.1 Morphological features

These features capture geometric aspects of the biofilm components.

“Biovolume” is the most commonly used feature in the biofilm community, where the number of pixels belonging to the biofilm is counted and then multiplied by the voxel size (here $0.13 \times 0.13 \times 0.5 \mu\text{m}$). A threshold is set before counting the pixels and only pixels above this threshold are counted. We generate multiple intensity thresholds, resulting in many biovolume features

that differ only by the threshold, so the algorithm must capture the values relevant for classification. This feature is calculated for the image as a whole and not per layers.

The “count” features give the number of bacteria on each layer. For this, a threshold is applied and then the biofilm components (here bacteria) are segmented using the connected components approach from Scipy Multidimensional Image Processing (`scipy.ndimage`). The objects are then counted for each layer and the counts per layer are aggregated to obtain one value per image. We generate multiple intensity thresholds, resulting in many count features that differ only by the threshold value, so that the algorithm can capture the relevant values for classification. The normalized values of the counts are calculated as described earlier. Although counts intuitively reflect the amount of bacteria, they are based on thresholding and only represent bacteria detected above that threshold, which may not capture all bacteria optimally.

Substratum Relative Coverage: Percentage of surface covered by biofilm in the 1st layer (surface of the well) of the image.

9.2 The statistics-based intensity features

The feature “globalMean” is extracted from the entire 3D image by calculating the average value of all pixel intensity values in the 3D image and is not derived layer by layer (Figure 3).

A set of intensity-based features labeled “maxdiffs”, “mindiffs”, “mdiffs” is calculated from the differences of the average pixel intensity values between two consecutive layers, i.e. the average value of the first layer is subtracted from the average value of the second layer and so on until 19 differences are obtained (Figure 3). Then the maximum, minimum and medium differences are selected and labeled as “maxdiffs”, “mindiffs”, “mdiffs” and given as individual values for each image.

The features “diff” and “diffNormalized” are calculated by subtracting the maximum pixel intensity value from the minimum pixel value for each layer separately. Consequently, we obtain several values for a feature (one per layer), which we combine into a single value by aggregation. In addition, we also introduce normalized features that are calculated for each layer, where the values are normalized by dividing the feature value for the individual




Figure 3: Schematic representation of calculation of intensity-based features.

layer by the sum of the values calculated for that layer. For example, the maximum difference in pixel intensity on layer 1 is divided by the sum of the maximum differences calculated for layers 1, 2, 3, 4 up to 21. Normalized features are indicated by the presence of “Normalized” in the feature name, e.g. “diffNormalized”.

Another set of intensity-based features labeled “mean”, “med”, “max”, “min” and “std” are calculated by computing the mean, median, maximum, minimum or standard deviation of the pixel intensity values for each layer, and these values are then aggregated. Normalized values are also calculated for these features by calculating the proportion for each layer.

9.3 Texture features

Texture features provide a description of an image that contains information about the spatial distribution of pixel intensity values. “MinProp” is calculated by counting the number of pixels in a layer whose pixel intensity value is below the average pixel intensity value in that layer. This feature is calculated separately for each layer and then aggregated for the entire image as described above. Normalized values are available as “MinPropNormalized”. This feature describes the empty space of the image that is not occupied by bacteria.

The “Homogeneity” is calculated based on the co-occurrence matrix, which represents the distribution of changes in the pixel intensity values of neighboring pixels in the X, Y and Z directions. The formula for calculating homogeneity has already been described¹. It measures the similarity of image structures that are spatially close to each other: a higher homogeneity indicates a more homogeneous image structure. This feature is calculated for the image as a

whole.

"ThicknessThreshold" is calculated by tracking the pixel intensity value of each pixel from the first layer (substrate) in height (along the normal to the substrate) until the pixel intensity value reaches 0. The number of the layer where the pixel intensity value reaches 0 is marked, and then the average value of these layers (heights) is calculated. One value is calculated per image.

"roughnessThreshold" is calculated by measuring the height differences for consecutive pixels calculated for "ThicknessThreshold" and then calculating the average. It reflects the height differences of the biofilm.

"Spreading" refers to the spread of the biofilm volume in the 3D image and is calculated as described before². Three different values are calculated: horizontal spreading (in xy-direction), "Sxy", vertical spreading (in z-direction), "Sz", and total spreading (in xyz-space).

9.4 Large language model (LLM) based features (GPT features)

GPT volume refers to the volume of the image layer occupied by the biofilm, where the number of bacteria is multiplied by the voxel size and then aggregated per image.

GPT fractal dimension captures the complexity of the geometric patterns on each image layer and is then aggregated into a single value. "GPTContrast", "GPTDissimilarity", "GPTHomogeneity", "GPTEnergy", "GPTCorrelation" and "GPTASM" were calculated using the Graycomatrix method. In this method, a gray-level co-occurrence matrix (GLCM) is created by counting how often a pixel with intensity i occurs in a specific spatial relationship to a pixel with intensity j . The correlation measures, for example, how strongly pairs of grey levels vary together, which indicates predictable spatial patterns in the image.

GPT correlation: Similar to the Pearson correlation coefficient, the correlation value is calculated based on the correlations between the values in graycomatrix. If the correlations are high, this indicates that certain grey level pairs occur in a predictable manner.

GPT Contrast: The contrast highlights the difference between gray-level values. Larger intensity differences contribute more to the sum, so a higher

contrast value indicates greater variability in the intensities of neighbouring pixels.

GPT dissimilarity: Similar to contrast, but it is calculated from absolute differences rather than a squared difference. Dissimilarity also captures local variations in grey levels, but its contribution grows more linearly compared to the quadratic term of contrast.

GPT homogeneity: It is calculated in a very similar way to non-GPT homogeneity, but here it is first calculated per layer and then summarised per image by aggregation. However, a high homogeneity indicates a more uniform and less varied texture.

GPT energy: Energy is another way of expressing the uniformity or regularity of texture. It is often used because it is on a scale of [0,1] when P is normalised.

A List and descriptions of biofilm features calculated by the pipeline

* The image can be analysed as a raw image without adjustments or the pixel intensity values can be normalised, which is referred to as norminten. The pixel intensity values in the raw image are normalized per layer by subtracting the average pixel intensity value of the individual layer by the sum of average pixel intensity values for all layers to obtain values between 0 and 1.

** All features in the data file are aggregated by calculating the mean, median, maximum, minimum and quartiles 10, 25, 75 and 90, standard deviation and variance across the image layers. If the name of the feature contains "SUMMARY", this means that it is an aggregated value of the feature, and the aggregation type is written after the underscore at the end of the feature name `_mean`, `_med`, `_min`, `_max`, `_q10`, `_q25`, `_q75`, `_q90`, `_std`, `_var`.

Feature name in the manuscript	Feature name in the datafile provided by the pipeline	Input type	normalization	Calculation per _	Description of feature	Threshold
counts (threshold < n)*	counts(inten<threshold-SUMMARYCustomAlgos.tsv	raw image	no	layer	The number of bacteria is determined in each layer of the 3D image by first applying the threshold value and then segmenting the bacteria, followed by counting the bacteria. The raw image or an image whose intensity values have been normalized is used as input. For the "normalized" feature, the counts per layer are normalized as a proportion of the sum of all values * 100 (to obtain percentages) per value.	Range from 0.01-0.2 increments for 0.01
counts (normalized intensity, threshold < n)*	counts(Norminten<threshold-SUMMARYCustomAlgos.tsv	normalized	no	layer		Range from 0.01-0.2 increments for 0.01
counts (normalized per layer, threshold < n)*	counts(inten<thresholdNormalized-SUMMARYCustomAlgos.tsv	raw image	yes	layer		Range from 0.01-0.2 increments for 0.01
counts (normalized intensity, normalized per layer, threshold < n)*	counts(Norminten<thresholdNormalized-SUMMARYCustomAlgos.tsv	normalized	yes	layer		Range from 0.01-0.2 increments for 0.01
Biovolume	BioVolumeThr-SUMMARYCustomAlgos.tsv	raw image	no	image	The biofilm volume is calculated by counting the number of pixels belonging to the biofilm after thresholding and multiplying it by the voxel size, which represents the physical dimensions of the pixel size in height, width and length (in micrometers); this gives the total volume occupied by the biofilm in the 3D image. Only one volume value is given per image and the aggregated values remain consistent.	Range from 0.01-0.29 increments for 0.01
GPT volume	GPTVolume-SUMMARYCustomAlgos.tsv	raw image	no	layer	The GPT volume is similar to the biovolume described above, but the bacteria are first counted layer by layer after the thresholding, and then the values of these layer counts are aggregated.	Fixed value
Normalized GPT volume	GPTVolumeNormalized-SUMMARYCustomAlgos.tsv	raw image	yes	layer		Fixed value
Substratum relative coverage	SubstratumRelativeCoverage-SUMMARYCustomAlgos.tsv	raw image	no	image	Percentage of the surface of the well (1st layer of the image) covered with biofilm.	Fixed value

*Feature "counts": Counts are a unique feature that captures the variations in counts of bacteria within bacteria with changing the threshold to detect the bacteria. It is an ensemble of features for which we know that they are correlated among each other and are useful to be used in machine learning to capture all aspects of thresholding and counting.

Figure 4: Morphological features.

Feature name in the manuscript	Feature name in the datafile provided by the pipeline	Input type	normalization	Calculation per _	Description of feature	Threshold
global mean	globalMean-SUMMARYDiffGlobal.tsv	raw image	no	image	The average pixel intensity of the entire 3D image reflects the overall brightness of the objects. A single value is calculated for each image, whereby all aggregated values are identical.	no
Layer-wise global maximum**	maxdiffs-SUMMARYDiffGlobal.tsv	raw image	no	image	The average pixel intensities are calculated for each layer of the 3D image, and the difference between the average of each layer and the next is calculated. Only the maximum difference is given, which represents the largest change in brightness between the layers and extends the understanding of the global mean.	no
Layer-wise global minimum**	mindiffs-SUMMARYDiffGlobal.tsv	raw image	no	image	The difference between the average intensities of consecutive layers is calculated in a similar way, but only the minimum difference is given to highlight areas of little or no change, which may correspond to stable biofilm regions or the background.	no
Layer-wise global mean**	mdiffs-SUMMARYDiffGlobal.tsv	raw image	no	image	The mean difference between the intensities of successive layers is calculated and reported and reflects the overall uniformity or variability of brightness in the vertical dimension of the biofilm; smaller values indicate a more uniform intensity and texture.	no
/	eigen-SUMMARYDiffGlobal.tsv	raw image	no	/	A value representing the differences in the overall 3D image is noted but is not currently relevant.	no
Intensity differences**	diff-SUMMARYCustomAlgos.tsv	raw image	no	layer	The difference between maximum and minimum pixel intensity within each layer is calculated and aggregated, where: (a) The mean aggregation indicates the image contrast, (b) The median reflects the typical intensity range, (c) The standard deviation indicates the variability of the contrast between the layers, and (d) The quantiles describe the distribution skewness of the contrast between the layers.	no
Normalized intensity differences**	diffNormalized-SUMMARYCustomAlgos.tsv	raw image	yes	layer	The differences between the maximum and minimum intensities per layer are normalized as a percentage of the total max-min sum across all slices and then aggregated to characterize the relative contrast distribution.	no
Mean intensity	mean-SUMMARYCustomAlgos.tsv	raw image	no	layer	The mean pixel intensity is calculated for each layer and then averaged across the layers, representing the average brightness per layer; higher values indicate denser or brighter regions and vice versa.	no
Normalized mean intensity	meanNormalized-SUMMARYCustomAlgos.tsv	raw image	yes	layer	The mean intensity on each layer is normalized as a percentage of the total sum of mean intensities per image before averaging across layers.	no

Figure 5: Statistics-based intensity features (1 of 2).

Feature name in the manuscript	Feature name in the datafile provided by the pipeline	Input type	normalization	Calculation per _	Description of feature	Threshold
Median intensity	med-SUMMARYCustomAlgos.tsv	raw image	no	layer	The median pixel intensity is calculated for each layer, and then the median values are aggregated across all layers.	no
Normalized median intensity	medNormalized-SUMMARYCustomAlgos.tsv	raw image	yes	layer	The median intensity per layer is normalized as a percentage of the total sum of the median intensities per image and then aggregated across layers.	no
Maximum intensity	max-SUMMARYCustomAlgos.tsv	raw image	no	layer	The maximum pixel intensity is calculated per layer and aggregated, highlighting localized bright spots such as clumps or active cells.	no
Normalized maximum intensity	maxNormalized-SUMMARYCustomAlgos.tsv	raw image	yes	layer	The maximum intensity per layer is normalized as a percentage of the total maximum per image and then aggregated.	no
Minimum intensity	min-SUMMARYCustomAlgos.tsv	raw image	no	layer	The minimum pixel intensity per layer is calculated and aggregated to capture the darkest or least dense biofilm regions; in combination with other statistics, it shows biofilm gaps, heterogeneity and spatial structure.	no
Normalized minimum intensity	minNormalized-SUMMARYCustomAlgos.tsv	raw image	yes	layer	The minimum intensity per layer is normalized as a percentage of the total minimum sum per image before aggregation.	no
Standard deviation of intensity	std-SUMMARYCustomAlgos.tsv	raw image	no	layer	The standard deviation of pixel intensities per layer is calculated and aggregated; low values indicate uniform intensity (low contrast), while high values indicate high contrast and high variability.	no
Normalized standard deviation of intensity	stdNormalized-SUMMARYCustomAlgos.tsv	raw image	yes	layer	The standard deviation per layer is normalized as a percentage of the total standard deviation sum per image and then aggregated.	no

***Layer wise features:*

Layer wise features capture how much the intensity fluctuates between the layers and how features behave across the volume

Figure 6: Statistics-based intensity features (2 of 2).

Feature name in the manuscript	Feature name in the datafile provided by the pipeline	Input type	normalization	Calculation per	Description of feature	Threshold
Min Prop	minProp-SUMMARYCustomAlgos.tsv	raw image	no	layer	The number of pixels whose intensity is below the average pixel intensity is counted per image. Min prop stands for dark areas, e.g. empty areas that are not occupied by biofilm.	no
Normalized Min Prop	minPropNormalized-SUMMARYCustomAlgos.tsv	raw image	yes	layer		no
Homogeneity	Homogeneity-SUMMARYCustomAlgos.tsv	raw image	no	image	Homogeneity is calculated using the co-occurrence matrix, which captures the distribution of intensity changes between neighbouring pixels in the X, Y and Z directions (after Beyenal et al., 2004). It measures how similar neighbouring image structures are, with higher values indicating greater uniformity. One value is given per image, whereby the aggregated values are identical.	no
Thickness	ThicknessThreshold=0.01-SUMMARYCustomAlgos.tsv	raw image	no	image	Pixel intensities are tracked vertically from the substrate layer, recording for each pixel the height at which the intensity drops to zero; the average of these heights is reported as a single image value with identical aggregation.	Range from 0.01-0.29 increments for 0.01
Roughnes	RoughnessThreshold=threshold-SUMMARYCustomAlgos.tsv	raw image	no	image	The average height difference between neighbouring pixels at a given layer thickness is calculated and reflects the roughness of the biofilm surface. One value per image is given, with uniform aggregation.	Range from 0.01-0.29 increments for 0.01
Horizontal spreading	SpreadingHorizontal-SUMMARYCustomAlgos.tsv	raw image	no	image	Three distributions (Dx, Dy, Dz) of the x, y and z coordinates of the biofilm pixels are created, as described by Mueller et al. (2006), and the variances for each axis are calculated. Horizontal spreading is calculated by summing up variances in x and y. Vertical spreading (in z direction) equals the variance of z. Total spreading (in z direction) is calculated by summing up variances in x, y and z.	no
Total spreading	SpreadingTotal-SUMMARYCustomAlgos.tsv	raw image	no	image		no
Vertical spreading	SpreadingVertical-SUMMARYCustomAlgos.tsv	raw image	no	image		no
GPT fractal dimension	GPTFractalDim-SUMMARYCustomAlgos.tsv	raw image	no	layer	The fractal dimension reflects the complexity and irregularity of the surface or boundary of the biofilm and quantifies the roughness at different scales; higher values indicate complicated and branched structures rather than smooth shapes.	no
Normalized GPT fractal dimension	GPTFractalDimNormalized-SUMMARYCustomAlgos.tsv	raw image	yes	layer		no
GPT homogeneity	GPTHomogeneity-SUMMARYCustomAlgos.tsv	raw image	no	image	This feature is calculated similarly to non-GPT homogeneity, but is first calculated per layer and then aggregated per image; higher homogeneity indicates a more uniform and consistent texture.	no

Figure 7: Texture features.

Feature name in the manuscript	Feature name in the datafile provided by the pipeline	Input type	normalization	Calculation per _	Description of feature	Threshold
GPT contrast	GPTContrast-SUMMARYCustomAlgos.tsv	raw image	no	image	Contrast emphasizes the difference between gray-level values. Larger differences in intensity contribute more heavily to the sum, so a higher contrast value indicates greater variability in the intensities of neighbouring pixels.	no
GPT correlation	GPTCorrelation-SUMMARYCustomAlgos.tsv	raw image	no	image	The correlation value is computed in a way that resembles the Pearson correlation coefficient based on the correlations between values in graycomatrix. If correlations are high this indicates that certain gray-level pairs occur in predictable way.	no
GPT dissimilarity	GPTDissimilarity-SUMMARYCustomAlgos.tsv	raw image	no	image	Similar to contrast, but it is calculated from absolute differences instead of a squared difference. Dissimilarity also captures local variations in gray levels, but its contribution grows more linearly compared to contrast's squared term. Dissimilarity is opposite to homogeneity.	no
GPT energy	GPTEnergy-SUMMARYCustomAlgos.tsv	raw image	no	image	Energy is another way to express the uniformity or orderliness in the texture. It is often used because it is on a [0,1] scale if P is normalized.	no

Figure 8: Graycomatrix method features.